

# Parallel Conjugate Gradient for Dense Matrices

Federico Betti

*MATH-454, Parallel and High Performance Computing*

EPFL — January 12, 2023



# MPI – Hard scaling


## Profiling/Hard scaling analysis

- Matrix vector products takes majority of running time.
- Amdahl's Law:  $S(p) \leq \frac{1}{\alpha}$ .

N	1024	2048	4096	8192
$\alpha$	0.0247	0.0090	0.0050	0.0019

## Parallelization strategy

- Row-wise domain decomposition.
- Reduction on  $\alpha_k$  and  $r_k$  and collective gathering for solution vector  $p_k$ .
- Computation cost  $\mathcal{O}\left(\frac{N^2}{p}\right)$ .
- Communication cost  $\mathcal{O}(\beta \frac{p-1}{p} N + \alpha \log_2 p)$  [2].



## Hard scaling results

- Speedup suffers from communications and synchronization overhead [1].
- Increasing the problem size boosts performance and mitigates drop with network communications ( $p > 36$ ).


# MPI – Weak scaling

## Weak scaling analysis

- For weak scaling, we fix the number of iterations for a fair comparison.
- To keep the workload  $\frac{N^2}{p}$  constant, from a serial problem size  $N_0$ , we then take  $N(p) = N_0\sqrt{p}$ .
- $\beta(N)$  is always rather small.

## Weak scaling results

- Gustafson's Law is too optimistic for relatively small starting points.
- There is evidence for improvement when increasing  $N_0$ .



## MPI Conclusions

- Synchronization overhead deteriorates parallelization potential. Hybrid OpenMP-MPI may provide benefits.
- Improvement for massively large problems.


# CUDA

## One thread per row approach

- Low GPU occupancy.
- Block-wise problem is too big to use shared memory.

## Key observation


- Computing equivalently  $A^T x$  favours memory coalescing.



## Multiple threads per row approach

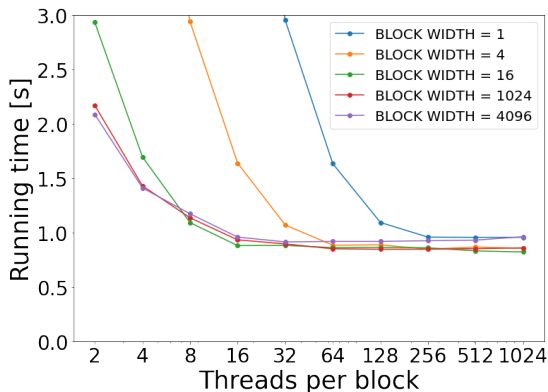
$$\begin{pmatrix} - & a_1 & - \\ - & a_2 & - \end{pmatrix} A = A^T \begin{pmatrix} | & | & | \\ a_1 & a_2 & | \\ | & | & | \end{pmatrix}$$

- Small block width gives queuing blocks. Large block width gives previous results.
- Can be improved using shared memory for  $x$ .



# References and appendix

- [1] R Oguz Selvitopi and Cevdet Aykanat. *Reducing Synchronization Overheads in CG-type Parallel Iterative Solvers by Embedding Point-to-point Communications into Reduction Operations.*
- [2] Rajeev Thakur, Rolf Rabenseifner, and William Gropp. *Optimization of collective communication operations in MPICH.*



N	$\alpha$
1024	0.0247
1448	0.0148
2048	0.0090
2896	0.0065
4096	0.0050
5792	0.0034
8192	0.0019
11585	0.0012
16384	0.0004